

# NEURAL MACHINE TRANSLATION (PART 1)

---

David Talbot

Spring 2017


Yandex School of Data Analysis

- Neural Networks (NN)
- Convolutional NN for NLP
- Recurrent Neural Networks and extensions (LSTM, GRU)
- First attempts at NMT (Recurrent Continuous Translation Model)
- Encoder-Decoders (sequence-to-sequence) models for MT


## PERCEPTRONS (MCCULLOUGH & PITTS, 1943)

Given input  $x = (x_1, x_2, \dots, x_m)$  where  $x_i \in \{0, 1\}$

$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$




What logic function does this compute?



## ADDITION

What weights and biases will implement  $f(x_1, x_2) = x_1 + x_2$ ?



Given training data

$$D = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$$

Choose parameters  $w$  and  $b$  such that  $\forall x_i \in D$

$$f(x_i) = \begin{cases} 1 & \text{if } w \cdot x_i + b > 0 \\ 0 & \text{otherwise} \end{cases} = y_i.$$

Perceptron learning rule

$$w_j \leftarrow w_j + \alpha(y_i - f(x_i))x_{ij}$$

Given training data

$$D = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$$


Measure the error on  $D$  using a cost-function, e.g.

$$C(w) = \frac{1}{2} \sum_{i=1}^n (y_i - f(x_i))^2$$

Minimize the error by updating  $w$  such that

$$w \leftarrow w - \alpha \nabla C(w)$$

## GRADIENT DESCENT






- $\alpha$  small: takes a long time to reach minimum of error function




- $\alpha$  large: may oscillate around minimum, without converging



## CONTINUOUS



## PERCEPTRON ACTIVATION FUNCTION



$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

## ADALINE LEARNING RULE (WIDROW, 1960)

Removing the step, makes the cost function differentiable


$$C(w) = \frac{1}{2} \sum_{i=1}^n (y_i - g(x_i))^2 \quad \text{where} \quad g(x_i) = w \cdot x_i + b$$

$$\begin{aligned} \frac{\partial C}{\partial w_j} &= \sum_{i=1}^n (y_i - w \cdot x_i - b) \frac{\partial}{\partial w_j} (y_i - w \cdot x_i - b) \\ &= - \sum_{i=1}^n (y_i - w \cdot x_i - b) x_{i,j} \end{aligned}$$

Gives us the Adaline update rule


$$w_j \leftarrow w_j + \alpha \sum_{i=1}^n (y_i - g(x_i)) x_{i,j}$$

## LINEARLY SEPARABLE CLASSES




Perceptron classifies blue as 0 and red as 1 with weights  
 $w = (?, ?)$ ,  $b = ?$

# MULTILAYER PERCEPTRON




$$g(x) = f^3\left(\sum_{k=1}^4 w_{1,k}^3 o_k + b\right)$$

# MULTILAYER PERCEPTRON



$$g(x) = f^3\left(\sum_{k=1}^4 W_{1,k}^3 f^2\left(\sum_{j=1}^3 W_{k,j}^2 o_j + b_k\right) + b\right)$$


# MULTILAYER PERCEPTRON



$$g(x) = f^3\left(\sum_{k=1}^4 w_{1,k}^3 f^2\left(\sum_{j=1}^3 w_{k,j}^2 f^1\left(\sum_{i=1}^5 (w_{j,i}^1 x_i + b_j)\right) + b_k\right) + b\right)$$




# MULTILAYER PERCEPTRON




What does this buy us if activations  $f^i(\cdot)$  are linear?

## SIGMOID



$$f(x) = \frac{1}{1 + \exp\{-w \cdot x + b\}}$$

# DEEP NETWORK



- Compute gradient on 'mini-batches' of the training data

$$\nabla C = \frac{\sum_{i=1}^n \nabla C_{x_i}}{n} \approx \frac{\sum_{j=1}^m \nabla C_{x_j}}{m}$$

- What assumptions do we need on our cost functions?

- Loss expressed as a function of the output layer
- Loss expressed as an average over data points

$$C_{mse} \equiv \frac{1}{2n} \sum_{i=1}^n \|y(x_i) - \hat{y}(x_i)\|^2$$

or

$$C_{cross\_entropy} \equiv -\frac{1}{n} \sum_{i=1}^n \sum_{y'} \Pr(y(x_i) = y') \log \Pr(\hat{y}(x_i) = y')$$

where  $y(x)$  and  $\hat{y}(x)$  are the true and predicted labels.

Compute derivatives for all parameters:

$$\frac{\partial C}{\partial w_{jk}^l}, \frac{\partial C}{\partial b_j^l} \quad \forall j, k, l$$

so that we can update the model to reduce the cost.

Recursion based on chain-rule: if  $f$  and  $g$  are both differentiable and  $h(x) = f(g(x))$  then

$$h'(x) = f'(g(x)) \cdot g'(x).$$

Let  $a^l$  be activation of  $j$ -th neuron at layer  $l$

$$a_j^l = \sigma\left(\sum_k w_{jk}^l a_k^{l-1} + b_j^l\right)$$

Using vectors

$$a^l = \sigma(w^l a^{l-1} + b^l)$$


Weighted input to  $j$ -th neuron at layer  $l$  (useful below)

$$z_j^l \equiv \sum_k w_{jk}^l a_k^{l-1} + b_j^l$$

Using vectors

$$z^l \equiv w^l a^{l-1} + b^l$$



# BACKPROPAGATION



\*Source: Michael Nielsen,  
<http://neuralnetworksanddeeplearning.com/chap2.html>




# BACKPROPAGATION




Can add  $\Delta z_j^l$  to input so output becomes  $\sigma(z_j^l + \Delta z_j^l)$ .

# BACKPROPAGATION




So overall cost changes by  $\frac{\partial C}{\partial z_j^l} \Delta z_j^l$

# BACKPROPAGATION



Define error at neuron as:  $\delta_j^l \equiv \frac{\partial C}{\partial z_j^l}$

# BACKPROPAGATION



If  $\frac{\partial C}{\partial z_j^l}$  is small, then not much to gain from changing  $z_j^l$

By applying chain-rule to the final layer

$$\delta_j^L \equiv \frac{\partial \mathcal{C}}{\partial z_j^L} \quad (1)$$

$$= \sum_k \frac{\partial \mathcal{C}}{\partial a_k^L} \frac{\partial a_k^L}{\partial z_j^L} \quad (2)$$

$$= \frac{\partial \mathcal{C}}{\partial a_j^L} \frac{\partial a_j^L}{\partial z_j^L} \quad (3)$$

$$= \frac{\partial \mathcal{C}}{\partial a_j^L} \sigma'(z_j^L) \quad (4)$$

Rate of change of  $C$  w.r.t. output activations

$$\delta^L = \nabla_{a^L} C \odot \sigma'(z^L)$$

In case of quadratic cost

$$\delta^L = (a^L - y) \odot \sigma'(z^L)$$

$\delta_l$  expressed in terms of  $\delta^{l+1}$

$$\begin{aligned}\delta_j^l &\equiv \frac{\partial \mathcal{C}}{\partial z_j^l} \\ &= \sum_k \frac{\partial \mathcal{C}}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{z_j^l} \\ &= \sum_k \frac{\partial z_k^{l+1}}{\partial z_j^l} \delta_k^{l+1}\end{aligned}$$

## BACKPROPAGATION: RECURSION (2)

Since

$$z_k^{l+1} = \sum_j w_{kj}^{l+1} a_j^l + b_k^{l+1} = \sum_j w_{kj}^{l+1} \sigma(z_j^l) + b_k^{l+1}$$

and differentiating

$$\frac{\partial z_k^{l+1}}{\partial z_j^l} = w_{kj}^{l+1} \sigma'(z_j^l)$$

so

$$\delta_j^l = \sum_k \frac{\partial z_k^{l+1}}{\partial z_j^l} \delta_k^{l+1} = \sum_k w_{kj}^{l+1} \delta_k^{l+1} \sigma'(z_j^l)$$



$$\delta^l = ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l)$$

$(w^{l+1})^T$  moves the error back through that layer

$\sigma'(z^l)$  moves it back through the activation


It turns out that

$$\frac{\partial \mathcal{C}}{\partial b_j^l} = \delta_j^l$$

So if the error is small, the bias will not change much


Similarly, it turns out that

$$\frac{\partial C}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$$




Weights from low activation neurons learn slowly

# HYPERBOLIC TANGENT




$$f(x) = \tanh(wx + b)$$

## RECTIFIED LINEAR UNIT (RELU)




$$f(x) = \max(0, wx + b)$$

## HIDDEN UNITS



A neural network with one hidden layer can approximate an arbitrary functions (with enough hidden units)

## FULLY CONNECTED NETWORKS



How about the inductive bias?

# CONVOLUTIONAL NETWORK

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature



# CONVOLUTIONAL NETWORK

1	1 <sub>x1</sub>	1 <sub>x0</sub>	0 <sub>x1</sub>	0
0	1 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	0
0	0 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	

Convolved  
Feature

## CONVOLUTIONAL NETWORK


1	1	1 <sub>x1</sub>	0 <sub>x0</sub>	0 <sub>x1</sub>
0	1	1 <sub>x0</sub>	1 <sub>x1</sub>	0 <sub>x0</sub>
0	0	1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved  
Feature

# CONVOLUTIONAL NETWORK



## WORD EMBEDDINGS: SPARSE VS DENSE REPRESENTATIONS

- Sparse: Each feature one dimension (binary value), each combination has its own dimension
- Dense: Each feature has a vector, no explicit encoding of feature combinations

# WORD EMBEDDINGS: SPARSE VS DENSE REPRESENTATIONS

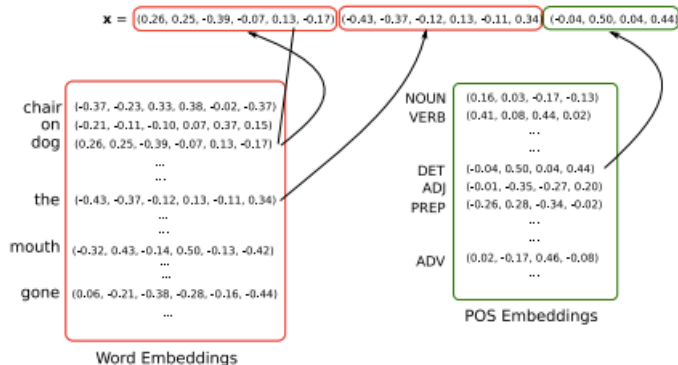
(a)

$\mathbf{x} = (0, \dots, 0, 1, 0, \dots, 0, 1, 0, \dots, 0, 1, 0, \dots, 0, 1, 0, \dots, 0, 0, 0, \dots, 0)$


Annotations above the vector  $\mathbf{x}$ :

- $w=\text{dog}$  points to the 4th element (1).
- $pw=\text{the}$  points to the 8th element (1).
- $pt=\text{NOUN}$  points to the 9th element (0).
- $pt=\text{DET}$  points to the 12th element (1).
- $w=\text{dog}\&pt=\text{DET}$  points to the 15th element (1).
- $w=\text{dog}\&pw=\text{the}$  points to the 18th element (1).
- $w=\text{chair}\&pt=\text{DET}$  points to the 21st element (0).

(b)




# WORD EMBEDDINGS



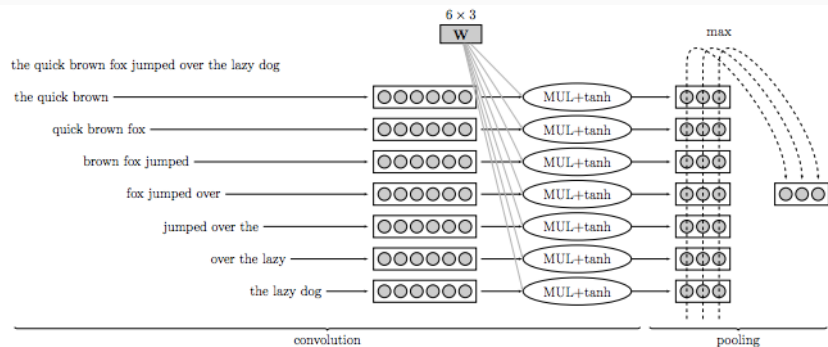
Skip gram model: predict word in random position close to  $w_t$

# WORD EMBEDDINGS



Magic of word embeddings? (More later)

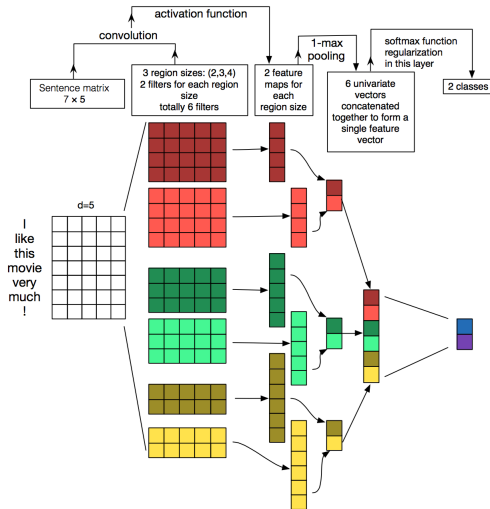
# CONVOLUTIONAL NETWORK



(Source: Goldberg, 2015)

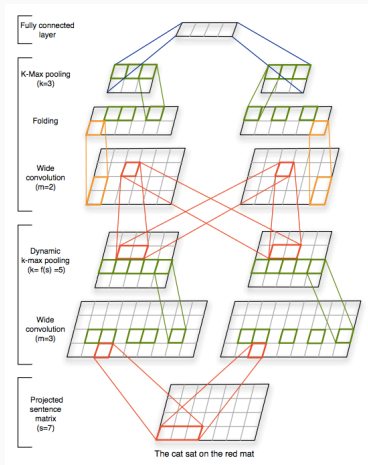


# CONVOLUTIONAL NETWORK



Source: Zhang, Y., & Wallace, B. (2015)


# CONVOLUTIONAL SENTENCE MODEL



Source: Kalchbrenner et al. (2015)

## CONVOLUTIONAL SENTENCE MODEL

*Encoder* in first NMT approach (Kalchbrenner & Blunsom 2013)



Source: Kalchbrenner et al. (2015)

Given training sequences of words  $w_1, \dots, w_T$  where  $w_t \in V$ , we want to learn a function


$$f(w_t, \dots, w_{t-n+1}) = \Pr(w_t | w_1^{t-1})$$

Bengio et al., 2003 decomposes  $f(\cdot)$  into

1. A mapping  $C$  from any element  $i$  of  $V$  to a real vector  $C(i) \in \mathbb{R}^m$  (a  $|V| \times m$  matrix)
2. A function (neural network) that assigns a probability  $P(w_t = i | w_1^{t-1})$  as

$$f(i, w_{t1}, \dots, w_{tn+1}) = g(i, C(w_{t1}), \dots, C(w_{tn+1}))$$

# NEURAL PROBABILISTIC LANGUAGE MODEL (BENGIO ET AL. 2003)



The output softmax layer is most computational

$$\Pr(w_t | w_1, \dots, w_{t-1}) = \frac{e^{y_{w_t}}}{\sum_{i \in V} e^{y_i}}$$

where


$$y = b + Wx + U \tanh(d + Hx)$$

and

$$x = (C(w_{t-1}), \dots, C(w_{t-n+1}))$$

- $W$  connects inputs to output directly (may be zero)
- $U$  connects hidden layer to output ( $|V| \times h$  matrix)
- $H$  connects inputs to hidden layer ( $h \times (n-1)m$  matrix)
- $b$  are input biases,  $d$  are hidden layer biases

- Number of parameters scales linearly with the vocabulary (unlike  $n$ -gram models)
- Embedding matrix  $C$  is shared among all inputs  $x_1, \dots, x_t$
- Main bottleneck is due to computation of softmax




State  $A_t$  updated from current input  $x_t$  and previous state  $A_{t-1}$

$$A_t = \tanh(Ux_t + WA_{t-1} + b) \quad \forall t \geq 1.$$

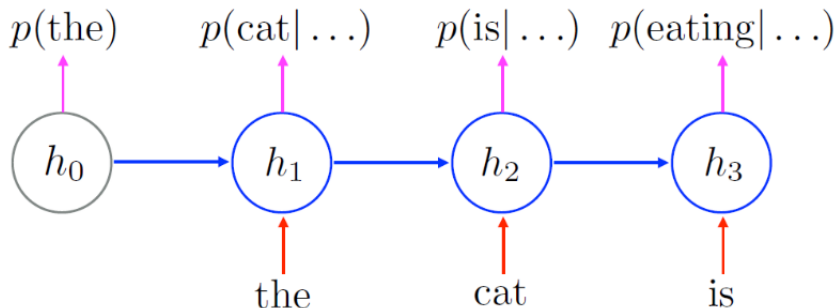


# RECURRENT NEURAL NETWORKS



Parameters shared across time steps

$p(\text{the, cat, is, eating})$




Kalchbrenner & Blunsom, (2013)

- First end-to-end NMT system
- Lower perplexity (average likelihood) than IBM models
- Encode source sentence with convolutional network
- RNN decoder generates target sentence conditioned on source sentence encoded in a ConvNN


$$\Pr(f|e) = \prod_{j=1}^J \Pr(f_j|f_{1:j-1}, e).$$

(See Appendix for details)

## RECURRENT CONTINUOUS TRANSLATION MODELS




# RECURRENT NEURAL NETWORKS




Why might backpropagation fail on such a 'deep' network?


# RECURRENT NEURAL NETWORKS



# LONG SHORT TERM MEMORY UNITS (HOCHREITER & SCHMIDHUBER, 1997)




Source: <http://colah.github.io/posts/2015-08-Understanding-LSTMs>






Separate cell  $C_t$  at each time frame to propagate information



## LSTM: FORGET GATE ACTIVATION


Controls how much each dimension of  $C_{t-1}$  is propagated to  $C_t$



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

## LSTM: NEW CANDIDATE CELL STATE


Constructs a preliminary cell state  $\tilde{C}_t$



$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

## LSTM: INPUT GATE ACTIVATION


Determines how much each dimension should be updated



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$


## LSTM: UPDATING THE CELL STATE

Forgetting some of  $C_{t-1}$  and overwriting with some of  $\tilde{C}_t$




$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

## LSTM: OUTPUT GATE




$$o_t = \sigma(W_o[h_{t-1}, x_t] + b_o)$$

$$h_t = o_t \odot \tanh(C_t)$$



## GATED RECURSIVE UNITS (CHO ET AL. 2014)



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$


$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t \odot h_{t-1}, x_t])$$

$$h_t = (1 - z_t) \odot h_{t-1} + z_t \odot \tilde{h}_t$$



# SEQUENCE TO SEQUENCE LEARNING (SUTSKEVER ET AL. 2014)



Deep LSTM encoder-decoder

- Encode source sentence with deep LSTM
- Generate target words from decoder LSTM after <EOS>
- Bootstrap training by reversing the source sentence (why?)

- Performance on long sentences is poor (why?)
- Open vocabulary translation is computationally hard
- Open vocabulary translation requires lots of data
- How can we make use of monolingual training data?

- Michael Nielson,  
<http://neuralnetworksanddeeplearning.com>
- C. Olah,  
<http://colah.github.io/posts/2015-08-Understanding-LSTMs>


## APPENDIX: RECURRENT CONTINUOUS TRANSLATION MODELS

---

- First end-to-end NMT system (Kalchbrenner & Blunsom, (2013))
- Lower perplexity (average likelihood) than IBM models
- Encode source sentence with convolutional network
- RNN decoder generates target sentence conditioned on source sentence encoded in a ConvNN

$$\Pr(f|e) = \prod_{j=1}^J \Pr(f_j|f_{1:j-1}, e).$$

## RECURRENT CONTINUOUS TRANSLATION MODELS



- $I \in \mathbb{R}^{q \times |V|}$  (input vocabulary embedding)
- $R \in \mathbb{R}^{q \times q}$  (recurrent transformation)
- $O \in \mathbb{R}^{|V| \times q}$  (output vocabulary mapping)

With  $v(f_i)$  as the one-hot representation of  $f_i \in V$

$$h_1 = \sigma(I \cdot v(f_1))$$



$$h_{i+1} = \sigma(R \cdot h_i + I \cdot v(f_{i+1}))$$

$$o_{i+1} = O \cdot h_i$$

$$\Pr(f_i = v | f_{1:i-1}) = \frac{\exp(o_{i,v})}{\sum_{v'=1}^{|V|} \exp(o_{i,v'})}$$



# RECURRENT CONTINUOUS TRANSLATION MODELS



Given a source sentence  $e = e_1, \dots, e_k$

- $E^e \in \mathbb{R}^{q \times k}$  (source sentence matrix)
- $K^i \in \mathbb{R}^{q \times i}$  ( $r = \lceil \sqrt{2N} \rceil$  convolution filters;  $N$  is max length)


Each convolution reduces dimensionality by  $i - 1$

$$\begin{aligned} E_1^e &= E^e \\ E_{i+1}^e &= \sigma(K^{i+1} * E_i^e) \end{aligned}$$


Result is a vector  $e \in \mathbb{R}^{q \times 1}$

( $L^j$  used to map convolved matrix  $E_i^e$  to  $\mathbb{R}^{q \times 1}$  if fewer than  $i + 1$  columns.)

# RECURRENT CONTINUOUS TRANSLATION MODELS




RCTM I



RCTM II

# RECURRENT CONTINUOUS TRANSLATION MODEL 1



- $I \in \mathbb{R}^{q \times |V^f|}$  (target vocabulary embedding)
- $R \in \mathbb{R}^{q \times q}$  (recurrent transformation)
- $O \in \mathbb{R}^{|V^f| \times q}$  (output vocabulary mapping)
- $S \in \mathbb{R}^{q \times q}$  (sentence mapping)

With  $csm(e)$  as the convolutional sentence model

$$s = S \cdot csm(e)$$

$$h_1 = \sigma(I \cdot v(f_1))$$

$$h_{i+1} = \sigma(R \cdot h_i + I \cdot v(f_{i+1}))$$

$$o_{i+1} = O \cdot h_i$$

$$\Pr(f_i = v | f_{1:i-1}, e) = \frac{\exp(o_{i,v})}{\sum_{v'=1}^{|V|} \exp(o_{i,v'})}$$

### Problems with Model 1

- RNN controls length of sentence (bias towards shorter sentences)
- Each target word is equally conditioned on the full source sentence


Model sentence length explicitly

$$\begin{aligned}\Pr(f|e) &= \Pr(m|e) \cdot \Pr(f|m, e) \\ &= \Pr(m|e) \prod_{i=1}^m \Pr(f_{i+1}|f_{1:i}, m, e)\end{aligned}$$

Where given source sentence  $e$  with length  $k$

$$\Pr(m|e) \approx \Pr(m|k) \approx \text{Poisson}(\lambda_k)$$

# RECURRENT CONTINUOUS TRANSLATION MODEL 2





Retain some positional information (cf. attention later)

- $E_i^e$  columns represent  $n$ -grams, e.g.  $E_2^e$  are bigrams
- Truncate convolutions to get 4-gram source features
- Invert convolutional sentence model on target side
- Target words are not uniformly conditioned on  $e$
- Additional parameters:  $T^{q \times q}$  translation matrix and inverted filters

- Trained with BPTT to optimize cross-entropy of parallel data
- Factor output to reduce computation

$$\Pr(f_i|f_{1:i-1}, e) = \Pr(C(f_i)|f_{1:i-1}, e) \cdot \Pr(f_i|C(f_i), f_{1:i-1}, e)$$

- Softmax with  $|V^f|$  terms reduces two with  $|C| + |V^f|/|C|$  terms

$$\Pr(C(f_i) = c|f_{1:i-1}, e) = \frac{\exp(o_{i,c})}{\sum_{c'=1}^{|C|} \exp(o_{i,c'})}$$

$$\Pr(f_i = v|(C(f_i), f_{1:i-1}, e) = \frac{\exp(o_{i,v})}{\sum_{v'=1}^{|C(f_i)|} \exp(o_{i,v'})}$$

WMT-NT	<i>2009</i>	<i>2010</i>	<i>2011</i>	<i>2012</i>
KN-5	218	213	222	225
RLM	178	169	178	181
IBM 1	207	200	188	197
FA-IBM 2	153	146	135	144
RCTM I	143	134	140	142
RCTM II	<b>86</b>	<b>77</b>	<b>76</b>	<b>77</b>

Table 1: Perplexity results on the WMT-NT sets.

WMT-NT PERM	<i>2009</i>	<i>2010</i>	<i>2011</i>	<i>2012</i>
RCTM II	174	168	175	178

Table 2: Perplexity results of the RCTM II on the WMT-NT sets where the words in the English source sentences are randomly permuted.

WMT-NT	<i>2009</i>	<i>2010</i>	<i>2011</i>	<i>2012</i>
RCTM I + WP	19.7	21.1	22.5	21.5
RCTM II + WP	19.8	21.1	22.5	21.7
cdec (12 features)	19.9	21.2	22.6	21.8

Table 4: Bleu scores on the WMT-NT sets of each RCTM linearly interpolated with a word penalty WP. The cdec system includes WP as well as five translation models and two language modelling features, among others.